

Indian Statistical Institute

Mid Semestral Examination: Machine Learning II

Maximum Marks: 60, Time: 2 Hrs. 15 Mins.

September 10, 2025

Answer, ALL questions. The maximum score you can obtain is 60.

- Let $p(k)$ be a one-dimensional discrete distribution that we wish to approximate, with support on nonnegative integers. One way to fit an approximating distribution $q(k)$ is to minimize the Kullback-Leibler (KL) divergence $KL(p||q)$. Show that when $q(k)$ is a Poisson distribution, this KL divergence is minimized by setting λ to the mean of $p(k)$. [10]
- Let $\{x^{(i)}\}_{i=1}^N$ be data samples from an unknown distribution. A Variational Autoencoder (VAE) models latent variables $z \in \mathbb{R}^k$ with prior

$$p(z) = \mathcal{N}(z; 0, I),$$

and likelihood

$$p_\theta(x|z) = \mathcal{N}(x; f_\theta(z), \sigma^2 I),$$

where f_θ is a neural network (decoder). Since the true posterior $p_\theta(z|x)$ is intractable, we introduce a variational approximation

$$q_\phi(z|x) = \mathcal{N}(z; \mu_\phi(x), \text{diag}(\sigma_\phi^2(x))),$$

with parameters ϕ (encoder network).

- Show that the marginal log-likelihood of a data point x can be decomposed as

$$\log p_\theta(x) = \mathbb{E}_{q_\phi(z|x)} \left[\log \frac{p_\theta(x, z)}{q_\phi(z|x)} \right] + \text{KL}(q_\phi(z|x) \parallel p_\theta(z|x)).$$

Conclude that maximizing the expected log-likelihood is equivalent to maximizing the Evidence Lower Bound (ELBO):

$$\mathcal{L}(\theta, \phi; x) := \mathbb{E}_{q_\phi(z|x)} [\log p_\theta(x|z)] - \text{KL}(q_\phi(z|x) \parallel p(z)).$$

- Prove that the ELBO is indeed a lower bound:

$$\mathcal{L}(\theta, \phi; x) \leq \log p_\theta(x),$$

and equality holds if and only if $q_\phi(z|x) = p_\theta(z|x)$ almost everywhere.

- Using the *reparameterization trick* $z = \mu_\phi(x) + \sigma_\phi(x) \odot \epsilon$, $\epsilon \sim \mathcal{N}(0, I)$, prove that the gradient of the ELBO w.r.t. θ and ϕ can be written as an expectation with respect to ϵ , thus avoiding direct differentiation through the sampling step.

[5 + 6 + 7 = 18]

3. Consider a very simple RNN with one hidden unit and scalar input defined by

$$h_t = \tanh(w_h h_{t-1} + w_x x_t), \quad y_t = w_y h_t,$$

where w_h, w_x, w_y are scalar weights, $h_0 = 0$, and $\tanh(z)$ is the hyperbolic tangent activation. Suppose the parameters are

$$w_h = 0.5, \quad w_x = 1.0, \quad w_y = 2.0,$$

and the input sequence is

$$x_1 = 1, \quad x_2 = -1, \quad x_3 = 2.$$

The target at the final time step is $y^* = 1$ and the loss is

$$L = \frac{1}{2}(y_3 - y^*)^2.$$

Compute the partial derivatives

$$\frac{\partial L}{\partial w_h} \quad \text{and} \quad \frac{\partial L}{\partial w_x},$$

showing the intermediate backward-pass quantities (e.g. $\delta_t = \partial L / \partial a_t$ with $a_t = w_h h_{t-1} + w_x x_t$) and the final numerical values. [6 + 6 = 12]

4. Consider the forward (noising) process used in denoising diffusion probabilistic models (DDPM). Let $x_0 \in \mathbb{R}^d$ be a data point, and let a fixed sequence $\{\beta_t\}_{t=1}^T$ satisfy $0 < \beta_t < 1$. Define

$$\alpha_t := 1 - \beta_t, \quad \bar{\alpha}_t := \prod_{s=1}^t \alpha_s.$$

The forward diffusion (a Markov chain) is defined by

$$q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{\alpha_t} x_{t-1}, \beta_t I), \quad t = 1, \dots, T.$$

- (a) Show that the marginal distribution $q(x_t | x_0)$ has the closed form

$$q(x_t | x_0) = \mathcal{N}(x_t; \sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t) I).$$

- (b) Using the result of (a), prove that x_t can be expressed explicitly as

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, \quad \epsilon \sim \mathcal{N}(0, I).$$

- (c) Show that the conditional posterior distribution is Gaussian:

$$q(x_{t-1} | x_t, x_0) = \mathcal{N}(x_{t-1}; \mu_q(x_t, x_0), \tilde{\beta}_t I),$$

where

$$\mu_q(x_t, x_0) = \frac{\beta_t \sqrt{\bar{\alpha}_{t-1}}}{1 - \bar{\alpha}_t} x_0 + \frac{\sqrt{\alpha_t} (1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t} x_t,$$

and

$$\tilde{\beta}_t = \frac{\beta_t (1 - \bar{\alpha}_{t-1})}{1 - \bar{\alpha}_t}.$$

[5 + 7 + 8 = 20]

5. Let us consider a Generative Adversarial Network (GAN) consisting of a Generator G and a Discriminator D . Let us also take the original data distribution as p_{data} , generated by G as p_g , and that of the noise as p_z . With these settings, if we replace the traditional Binary Cross Entropy loss in GANs with the Mean Squared Error loss then the modified objective function becomes as follows:

$$\begin{aligned} \min_D &= \frac{1}{2} \mathbb{E}_{x \sim p_{data}} (D(x) - b)^2 + \frac{1}{2} \mathbb{E}_{z \sim p_z} (D(G(z)) - a)^2, \\ \min_G &= \frac{1}{2} \mathbb{E}_{z \sim p_z} (D(G(z)) - c)^2. \end{aligned}$$

Let us call this an MSE (Mean Squared Error)-GAN, where a , b , and c , respectively denote the generated data label, the real data label, and the label of the data that the generator wants the discriminator to believe.

- What can be an intuition behind replacing the Binary Cross Entropy loss with a Mean Squared Error?
- There exists a family of divergence measures called f -divergence. The general form of f divergence between a couple of distributions P and Q are defined as:

$$D(P||Q) = \int f\left(\frac{p(x)}{q(x)}\right) q(x) dx.$$

Show that the KL divergence is actually a special case of f -divergence for a particular form of the function f , and also derive the form of this generating function.

- The χ^2 divergence is defined as $\chi^2(P||Q) = \int \frac{(p(x)-q(x))^2}{q(x)} dx$. Show that this divergence is a member of the f -divergence family and also derive the corresponding generating function f .
- If we impose the constraints $b - c = 1$ and $b - a = 2$ then can you show that the objective function of MSE-GAN can also be reduced to a χ^2 divergence between $p_{data} + p_g$ and $2p_g$ in a manner similar to the canonical GAN? [3 + 4 + 6 + 7 = 20]

Indian Statistical Institute

End Semester Examination

M. Tech. (CS), Machine Learning II

Maximum Marks: 100

Duration: 3 Hrs. 15 Mins.

Note:

- Attempt all questions.
 - Assume suitable values wherever necessary and state assumptions clearly.
 - All notations follow standard definitions used in deep learning and transformer architectures.
 - Use of calculator is permitted.
-

1. Consider a linear autoencoder defined as follows:

$$x \in \mathcal{R}^n, \quad h = W_1 x, \quad \hat{x} = W_2 h,$$

where $W_1 \in \mathcal{R}^{m \times n}$ is the encoder weight matrix and $W_2 \in \mathcal{R}^{n \times m}$ is the decoder weight matrix.

- If the hidden layer dimension $m = n$, determine what the product $W_2 W_1$ must be for perfect reconstruction of all inputs x .
- If $m < n$, explain mathematically why perfect reconstruction of all inputs is not possible, and describe what $W_2 W_1$ represents in this case.
- Briefly interpret the geometric meaning of part (b): what kind of subspace does the autoencoder learn when $m < n$?

[3 + 3 + 2 = 8]

2. Let $p(k)$ be a one-dimensional discrete probability distribution defined on nonnegative integers $k = 0, 1, 2, \dots$, which we wish to approximate by a Poisson distribution $q(k; \lambda)$ having parameter $\lambda > 0$. The Poisson distribution is given by

$$q(k; \lambda) = \frac{e^{-\lambda} \lambda^k}{k!}, \quad k = 0, 1, 2, \dots$$

- Show that minimizing the KL divergence $KL(p \| q_\lambda)$ with respect to λ leads to

$$\lambda^* = \sum_{k=0}^{\infty} k p(k),$$

that is, the optimal Poisson parameter equals the mean of $p(k)$. Provide all intermediate steps clearly.

- Explain intuitively why this result implies that the “best” Poisson approximation to any discrete nonnegative distribution matches its expected value.

- (c) Show that $KL(p\|q_\lambda)$ is a convex function of λ , and hence verify that the stationary point $\lambda^* = E_p[k]$ corresponds to a unique global minimum.

$$[6 + 3 + 3 = 12]$$

3. (a) A simple Recurrent Neural Network (RNN) operates on an input sequence

$$x = [x_1, x_2] = [1, 2],$$

with parameters:

$$W_x = 1.0, \quad W_h = 0.5, \quad b = 0, \quad \text{and activation } f(z) = \tanh(z).$$

Assume the initial hidden state $h_0 = 0$.

- i. Compute h_1 and h_2 step-by-step, showing all intermediate values before and after applying \tanh .
- ii. If the output is computed as $y_t = h_t$, write the final output sequence $[y_1, y_2]$.

- (b) Consider an LSTM cell with the same input sequence $x = [1, 2]$ and the following parameters:

$$W_i = W_f = W_o = W_c = 1, \quad U_i = U_f = U_o = U_c = 0, \quad b_i = b_f = b_o = b_c = 0,$$

and the initial states $h_0 = 0, c_0 = 0$. Use the activation functions $\sigma(z) = \frac{1}{1+e^{-z}}$ and $\tanh(z)$.

- i. Compute the gate activations (i_t, f_t, o_t) and candidate cell state (\tilde{c}_t) for $t = 1, 2$.
- ii. Compute the cell states c_1, c_2 and hidden states h_1, h_2 .
- iii. Compare how the LSTM's memory cell c_t differs from the RNN's hidden state in preserving information across time steps.

$$[(4+4)+(4+5+3) = 20]$$

4. A Transformer's Scaled Dot-Product Attention mechanism operates using three matrices: Query (Q), Key (K), and Value (V). You are given the following input matrices for a single attention head:

$$Q = \begin{bmatrix} 1 & 0 & 1 \end{bmatrix}, \quad K = \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 1 & 0 \end{bmatrix}, \quad V = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}$$

- (a) Compute the unnormalized attention scores.
- (b) Apply the appropriate scaling factor to the computed scores to obtain the scaled attention scores.
- (c) Calculate the normalized attention weights using the Softmax function applied to the scaled scores.
- (d) Using the obtained attention weights, compute the final attention output vector as a weighted sum of the value vectors.

$$[4 + 5 + 4 + 7 = 20]$$

5. Explainable AI techniques such as LIME and SHAP help interpret the contribution of individual input features to a model's prediction. Consider a simple linear regression model used for predicting a student's exam score based on two features:

$$f(x_1, x_2) = 5 + 3x_1 + 2x_2$$

where x_1 = hours studied and x_2 = number of practice tests taken. We wish to explain the prediction for a specific student with input $(x_1, x_2) = (2, 1)$.

- (a) Using the SHAP framework, assume the following coalition function values (model outputs when only a subset of features is present):

$$\begin{aligned} f(\emptyset) &= 5, \\ f(\{x_1\}) &= 11, \\ f(\{x_2\}) &= 9, \\ f(\{x_1, x_2\}) &= 15. \end{aligned}$$

Compute the Shapley values ϕ_{x_1} and ϕ_{x_2} for this prediction and verify that:

$$f(x_1, x_2) = f(\emptyset) + \phi_{x_1} + \phi_{x_2}.$$

- (b) Interpret the computed Shapley values: Which feature contributes more to the model's output for this particular prediction, and by how much relative to the baseline?

$$[5 + 3 = 8]$$

6. Let $s \in R^n$ be a vector of attention scores with components s_j , and define the *softmax* function

$$a = \text{softmax}(s), \quad a_i = \frac{e^{s_i}}{\sum_{t=1}^n e^{s_t}}, \quad i = 1, \dots, n.$$

The vector a represents the attention weights corresponding to s .

- (a) Prove that the softmax function is invariant to adding the same constant to all entries of s ; that is, for any scalar $c \in R$,

$$\text{softmax}(s + c\mathbf{1}) = \text{softmax}(s),$$

where $\mathbf{1}$ denotes the all-ones vector.

- (b) Derive the *Jacobian matrix* of the softmax function, defined as

$$J_s = \frac{\partial a}{\partial s} \in R^{n \times n}, \quad \text{where} \quad (J_s)_{ij} = \frac{\partial a_i}{\partial s_j}.$$

Show that $J_s = \text{diag}(a) - aa^\top$ and hence that for all i, j ,

$$\frac{\partial a_i}{\partial s_j} = a_i(\delta_{ij} - a_j),$$

where δ_{ij} is the Kronecker delta.

(Hint: Use the quotient rule starting from $a_i = e^{s_i} / \sum_t e^{s_t}$.)

[4+6=10]

7. (a) Consider a graph G with node indices $\{0, 1, 2, 3, 4\}$ and edge indices given by:

$$\begin{bmatrix} 0 & 0 & 0 & 1 & 2 & 2 & 3 & 3 & 3 & 4 \\ 1 & 2 & 3 & 0 & 0 & 3 & 0 & 2 & 4 & 3 \end{bmatrix}$$

Draw the computation graph for node ID 2 for a 2-hop message passing operation. Assume the node feature matrix of G is:

$$X = \begin{bmatrix} 0 & -1 \\ 1 & -2 \\ 0 & 2 \\ 1 & 0 \\ -1 & 3 \end{bmatrix}$$

Estimate the aggregated features for node 2 using a bottom-to-top approach with the mean aggregator (no self-loop addition required). Further, compute A^2X , where A is the normalized adjacency matrix, to aggregate two-hop node features. Show that the updated node features for node 2 are equal under both aggregation approaches.

- (b) Two nodes p and q in a connected graph respectively have degrees m and n . Assume all nodes possess identical features. If neighborhood aggregation is performed, design two different aggregators that yield distinct neighborhood representations for nodes p and q under the following cases:
 - (a) $m = n$
 - (b) $m \neq n$
- (c) Let L be a graph Laplacian with two eigenvectors

$$u = [-1 \quad 3a \quad b], \quad v = [2b \quad c \quad -a]$$

corresponding to eigenvalues greater than zero, where $a, b, c \in R$. Find the values of a, b , and c that satisfy orthogonality and normalization conditions of Laplacian eigenvectors.

- (d) Discuss three advantages of employing **Knowledge Distillation** from a larger (teacher) network to a smaller (student) network, in comparison to designing an end-to-end trainable compact model.

[8 + 4 + 5 + 3 = 20]

8. (a) A binary classifier predicts whether a candidate is *selected* ($\hat{Y} = 1$) or *not selected* ($\hat{Y} = 0$) in a job interview for the position of Data Scientist. The following table shows predictions for 12 candidates, equally divided between males and females.

Candidate ID	Gender	Predicted Outcome (\hat{Y})
1	Female	0
2	Male	1
3	Male	1
4	Female	1
5	Male	1
6	Female	0
7	Female	0
8	Male	0
9	Female	1
10	Male	1
11	Female	0
12	Male	1

- i. Compute the *Demographic Parity Difference (DPD)* between male and female candidates.
- ii. Interpret your result: which gender group receives more favorable predictions?
- (b) A probabilistic language model assigns conditional probabilities to each token in a sequence (w_1, w_2, \dots, w_N) . The perplexity (PP) of the model is defined as:
- $$PP = \exp \left(-\frac{1}{N} \sum_{i=1}^N \log P(w_i | w_1, \dots, w_{i-1}) \right).$$
- i. Derive this expression starting from the cross-entropy formulation
- $$H(p, q) = -\frac{1}{N} \sum_{i=1}^N \log q(w_i).$$
- ii. Compute the perplexity for a 4-token sequence where the model assigns probabilities $P(w_1) = 0.4$, $P(w_2|w_1) = 0.3$, $P(w_3|w_1, w_2) = 0.2$, and $P(w_4|w_1, w_2, w_3) = 0.1$.
- iii. Interpret what this perplexity value implies about the model's confidence and prediction quality.

$$[(4+2)+(2+2+2) = 12]$$